

ADVAIT JISHNANI

Portfolio Content Source of Truth

Software Engineering | Data Engineering | AI

Version: October 2026

Purpose. Use this document as the factual content authority for the personal portfolio website. Website copy may be shortened or paraphrased, but names, dates, technologies, metrics, project outcomes, and employment claims should not be changed or expanded beyond what is recorded here.

LOCATION	New York, NY
PRIMARY EMAIL	advaitnaresh@gmail.com
ACADEMIC EMAIL	aj4700@nyu.edu
LINKEDIN	linkedin.com/in/advait-jishnani
GITHUB	github.com/advaitnaresh
EDUCATION	M.S. Computer Science, New York University (2025-2027); B.E. Computer Science, Minor in Finance, BITS Pilani Goa (2020-2024)
CORE POSITIONING	Software engineering, data engineering, backend systems, distributed data platforms, and applied AI/LLM systems

[LinkedIn](#) | [GitHub](#)

1. Profile and education

PROFESSIONAL SUMMARY

Software and data engineer pursuing an M.S. in Computer Science at New York University. Experience spans production data and MLOps systems at Visa, creator-facing software and messaging infrastructure at Return on Creators, and student-facing workflow automation at NYU Tandon Digital Learning. Technical strengths include Python, SQL, Java, React, distributed data processing, cloud systems, vector search, and applied AI.

WEBSITE POSITIONING

Lead with the intersection of software engineering, data engineering, and AI. The strongest proof points are production ownership, measurable performance improvements, reliability work, full-stack feature delivery, and technically substantial personal projects.

Education

NEW YORK UNIVERSITY

Master of Science in Computer Science | New York, NY |
Sept 2025 - May 2027 | GPA: 4.0/4.0

BITS PILANI, GOA

Bachelor of Engineering in Computer Science, Minor in
Finance | Goa, India | Aug 2020 - May 2024 | GPA: 4.0/4.0

TECHNICAL FOCUS

Data engineering, distributed systems, cloud platforms, backend development, machine learning, LLM-powered retrieval, and reliable production software.

2. Professional experience

New York University - Tandon Digital Learning

Salesforce Assistant | Dec 2025 - Present | New York, NY

ORGANIZATION CONTEXT

NYU Tandon Digital Learning designs and supports flexible online and hybrid STEM education, including degree, professional, and executive programs, while partnering with faculty and staff on digital teaching and student services.

Public organization source: [NYU Tandon Digital Learning](#)

WHAT I DID

Build and maintain student-facing systems and operational workflows across Salesforce, Stripe, Brightspace, and custom web experiences. The work combines application development, integration engineering, workflow automation, and data-model design.

VERIFIED HIGHLIGHTS

- Rebuilt the Stripe payment and enrollment workflow with Salesforce Flow automations for multiple NYU Tandon courses supporting \$300K+ in revenue.
- Designed an Aura-based web application with reusable templates and embedded onboarding paths for student-facing workflows.
- Integrated Salesforce, Stripe, and Brightspace through automated workflows and REST APIs, reducing manual handoffs and improving cross-system reliability.

TECHNOLOGIES AND SYSTEMS

Salesforce, Salesforce Flow, Aura, Stripe, Brightspace, REST APIs, workflow automation, data modeling

2. Professional experience - Return on Creators

Return on Creators

Data Engineering Intern | Jun 2026 - Aug 2026 | New York, NY

ORGANIZATION CONTEXT

Return on Creators is a creator-economy company focused on creator commerce, influence, and performance measurement. Its platform and operating model evaluate creators against brands and connect creator activity to commercial outcomes.

Public organization source: [Return on Creators](#)

WHAT I DID

Worked across frontend, backend, messaging infrastructure, data modeling, and automation. The internship centered on production systems used to onboard and communicate with creators reliably at scale.

VERIFIED HIGHLIGHTS

- Built a full-stack React admin inbox with real-time messaging and underlying Python services, shipping production features across frontend, backend, and data layers.
- Shipped an end-to-end iMessage creator onboarding flow with Supabase and automated routing, designed to support up to 1,000 simultaneous creators while preserving production reliability.
- Developed a mass-messaging workflow designed to reach 500 creators without disrupting the existing backend, accounting for load, failure modes, and operational reliability.
- Re-keyed a production messaging table to a composite key, rebuilt inbound routing across multiple messaging lines, and resolved 3 data-integrity bugs through testing and production debugging.
- Built a persona-extraction workflow that turned creators' text and voice-call transcripts into structured profile data.
- Planned a personalized LLM extension so creators could receive contextual answers inside the onboarding flow; this was a planned extension rather than a shipped production claim.

TECHNOLOGIES AND SYSTEMS

React, Python, Supabase, SQL/data modeling, real-time messaging, APIs, workflow automation, LLM/AI integration planning

2. Professional experience - Visa

Visa Inc.

Data Engineer | Jan 2024 - Aug 2025 | Bengaluru, India

ORGANIZATION CONTEXT

Visa is a global digital payments technology network connecting consumers, businesses, financial institutions, and governments across more than 200 countries and territories.

Public organization source: [Visa - About Visa](#)

WHAT I DID

Built and operated production data and model infrastructure for large-scale financial workflows, with emphasis on automation, data quality, orchestration, observability, and performance optimization.

VERIFIED HIGHLIGHTS

- Automated the end-to-end Model Risk Management (MRM) process and model-refitting pipeline, cutting turnaround time by 60% and manual effort by 80%.
- Implemented automated testing, validation, monitoring, and alerting across 15+ production models to catch data-quality and model issues before downstream release.
- Orchestrated dependencies, retries, scheduling, and monitoring across 20+ Apache Airflow production workflows.
- Engineered large-scale financial-data transformations with Python, Java, SQL, PySpark, and Hive.
- Supported Tableau reporting and monitoring used by 30+ cross-functional stakeholders.

TECHNOLOGIES AND SYSTEMS

Python, Java, SQL, PySpark, Spark, Hive, Apache Airflow, Tableau, data quality, observability, workflow orchestration, MLOps

Visa Inc.

Data Engineer Intern | May 2023 - Jun 2023 | Bengaluru, India

ORGANIZATION CONTEXT

This internship was part of Visa's large-scale data engineering environment supporting production analytical workloads.

WHAT I DID

Focused on query and distributed-processing performance, investigating execution bottlenecks and applying Hive/Spark optimization techniques to production workloads.

VERIFIED HIGHLIGHTS

- Reduced a production query from 2 hours to under 12 minutes, a 90% improvement.
- Applied partitioning, caching, and bucketing in Hive and Spark to improve execution speed and scalability.
- Investigated backend bottlenecks and schema-design trade-offs with senior engineers.

TECHNOLOGIES AND SYSTEMS

Hive, Apache Spark, SQL, distributed data processing, performance tuning, schema design

3. Selected project catalog

These descriptions are intended to be safe for direct website use. For projects without a verified individual public repository in the current source set, the website should either link to the main GitHub profile or omit the code link until the exact repository URL is confirmed.

RoCathon: Hybrid Creator Search

AI / SEARCH / DATA ENGINEERING | Individual repository not yet verified - use GitHub profile or omit code link

WEBSITE-READY DESCRIPTION

A hybrid creator search system over 10,000+ profiles that combines structured metadata, lexical retrieval, vector search, and Gemini embeddings for higher-quality discovery.

DETAILED SOURCE DESCRIPTION

Built a semantic search layer using PostgreSQL and pgvector, combining lexical matching with HNSW vector retrieval across structured and semantic creator signals. Added metadata filters and a commerce-aware re-ranking stage so results could be ordered using more than embedding similarity alone.

VERIFIED FACTS AND METRICS

- Won 1st place at RoCathon.
- Indexed 10,000+ creator profiles.
- Improved top-result relevance by 35% through reranking, metadata filters, semantic retrieval, and Gemini embeddings.

TECHNOLOGY

PostgreSQL, pgvector, TypeScript, Node.js, Gemini API, embeddings, HNSW, hybrid search

Code / source: <https://github.com/advaitnaresh>

Aries: Real-Time Data Platform

DATA ENGINEERING / STREAMING | Individual repository not yet verified - use GitHub profile or omit code link

WEBSITE-READY DESCRIPTION

A real-time brand-intelligence data platform that processes high-volume streaming events and serves low-latency KPIs through a tiered storage architecture.

DETAILED SOURCE DESCRIPTION

Designed a streaming pipeline with Kafka and Spark for continuous event processing and downstream analytics. Used MinIO and PostgreSQL as a tiered storage layer to balance scalable event retention with fast access to live metrics. The current portfolio draft also describes a live dashboard and sentiment-analysis layer.

VERIFIED FACTS AND METRICS

- Processed 10,000+ streaming events per second.
- Served live KPIs at under 100 ms latency.

TECHNOLOGY

Apache Kafka, Apache Spark, Python, PostgreSQL, MinIO; current portfolio draft also references Streamlit and sentiment analysis

Code / source: <https://github.com/advaitnaresh>

3. Selected project catalog - continued

VulCAN: Automated Cloud Vulnerability Analyzer

CLOUD / SECURITY / BACKEND | Individual repository not yet verified - use GitHub profile or omit code link

WEBSITE-READY DESCRIPTION

An event-driven cloud security system that automates codebase vulnerability scanning and decouples high-volume remediation work through asynchronous AWS services.

DETAILED SOURCE DESCRIPTION

Built a cloud-native scanning service deployed with AWS CDK. Used SQS to decouple asynchronous jobs and Fargate for scalable processing, with Python and PostgreSQL supporting application logic and persistence.

VERIFIED FACTS AND METRICS

- Accelerated security-remediation workflows by 75%.
- Sustained 5,000+ daily jobs without queue bottlenecks.

TECHNOLOGY

AWS Fargate, AWS SQS, AWS CDK, Python, PostgreSQL

Code / source: <https://github.com/advaitnaresh>

Foreign Whispers

AI / MEDIA SYSTEMS | Current portfolio project - individual repository not yet verified

WEBSITE-READY DESCRIPTION

An AI media pipeline for transcription, translation, speech synthesis, and synchronized video output, with API orchestration separated from GPU-intensive inference.

DETAILED SOURCE DESCRIPTION

The project is structured as a media-processing pipeline rather than a monolithic AI demo: application/API orchestration is separated from model-heavy inference, while FFmpeg handles media synchronization and output assembly.

TECHNOLOGY

FastAPI, Whisper, Docker, FFmpeg, speech and media-processing APIs

AI-Powered Applicant Tracking System (shivarkATS)

AI / CLOUD APPLICATIONS | Verified public repository

WEBSITE-READY DESCRIPTION

A cloud-based applicant tracking system with an ML-powered resume-screening workflow designed to reduce manual recruiting effort.

DETAILED SOURCE DESCRIPTION

Architected and co-developed an ATS on AWS with machine-learning-assisted resume screening. The project demonstrates application design, cloud deployment, and applied ML in a recruiting workflow.

VERIFIED FACTS AND METRICS

- GitHub profile description states a 60% reduction in manual screening effort.

TECHNOLOGY

AWS, Python, MongoDB, Machine Learning, web frontend

Code / source: <https://github.com/advaitnaresh/shivarkATS>

3. Selected project catalog - continued

Model Predictive Control in Robotics

ROBOTICS / OPTIMIZATION | Verified public repository

WEBSITE-READY DESCRIPTION

A robotics control project exploring Model Predictive Control (MPC), control horizons, and computational trade-offs for choosing optimized actions toward a target state.

DETAILED SOURCE DESCRIPTION

Developed and analyzed linear and multivariable MPC formulations under Prof. Amit Setia. The work explored how reducing state/control variables and shortening the control horizon can reduce online computation while preserving useful control behavior. The repository includes MATLAB implementations and simulation material.

VERIFIED FACTS AND METRICS

- Repository contains 13 commits and multiple MPC implementations/simulation artifacts.
- Current portfolio material separately describes a TurtleBot 3 / ROS application; use the broader robotics description unless a dedicated ROS source is provided.

TECHNOLOGY

MATLAB, Model Predictive Control, optimization, robotics

Code / source: <https://github.com/advaitnaresh/MPC>

Machine Learning in Financial Markets

MACHINE LEARNING / TIME SERIES | Verified public repository

WEBSITE-READY DESCRIPTION

A financial time-series project comparing moving-average approaches with ARIMA forecasting on historical Apple stock data.

DETAILED SOURCE DESCRIPTION

Built a 5-day regressive moving-average baseline and then developed an ARIMA model using approximately 10 years of Apple stock data (2009-2019). The repository documents experimentation, model notebooks, and cloud-sourced data.

VERIFIED FACTS AND METRICS

- Repository README reports approximately 82% accuracy for the ARIMA model.
- Repository work was completed under Prof. J. K. Sahoo at BITS Pilani Goa.

TECHNOLOGY

Python, Jupyter, Pandas, ARIMA, GCP, time-series analysis

Code / source: <https://github.com/advaitnaresh/Machine-Learning-In-Financial-Markets>

3. Selected project catalog - continued

Restaurant Concierge Bot

CLOUD / SERVERLESS APPLICATIONS | Verified public repository

WEBSITE-READY DESCRIPTION

A serverless dining concierge chatbot that collects user preferences conversationally, searches restaurant data, and sends personalized recommendations by email.

DETAILED SOURCE DESCRIPTION

Built a microservice-driven AWS architecture with Lex for conversation handling, API Gateway and Lambda for request processing, SQS for asynchronous work, DynamoDB and OpenSearch for hybrid restaurant storage/search, and SES for email delivery. A Python ingestion script loads restaurant data from the Yelp API.

VERIFIED FACTS AND METRICS

- Ingests 1,000+ restaurant records from the Yelp API.
- Uses SQS to decouple user requests from background recommendation processing.

TECHNOLOGY

AWS Lambda, Amazon Lex, API Gateway, SQS, DynamoDB, OpenSearch, SES, Python, HTML/CSS/JavaScript

Code / source: <https://github.com/advaitnaresh/Restaurant-Concierge-Bot>

4. Skills and technical taxonomy

LANGUAGES	Python, SQL, Java, C++, TypeScript, JavaScript, Bash
SOFTWARE & BACKEND	React, Node.js, REST APIs, distributed systems, Git/GitHub, CI/CD
DATA ENGINEERING	PySpark, Spark, Airflow, Kafka, Hive, ETL/ELT, data modeling, data quality, query optimization
DATABASES & STORAGE	PostgreSQL, pgvector, MongoDB, Supabase, MinIO, DynamoDB, OpenSearch
CLOUD & INFRASTRUCTURE	AWS, GCP, Docker, Kubernetes, AWS CDK, Fargate, SQS, Lambda
AI & ANALYTICS	LLMs, RAG, embeddings, vector search, HNSW, Gemini API, machine learning, model evaluation, ARIMA, Tableau
ROBOTICS & MEDIA	Model Predictive Control, MATLAB, FFmpeg, Whisper

5. Website content rules

- The website may shorten descriptions, but it should not increase a metric, change a date, or imply a shipped feature when the source says planned.
- Use Software Engineer / Data Engineer / AI-focused Engineer as broad positioning language. Do not invent an official specialization that is not listed on the degree record.
- For Return on Creators, the personalized LLM onboarding assistant should be described as planned/design work unless a newer source confirms it shipped.
- For project code buttons, only use the individual GitHub URLs explicitly listed in this document. For RoCathon, Aries, VulCAN, and Foreign Whispers, either link to the GitHub profile or omit the code link until the exact repository is confirmed.
- For the MPC project, the verified public repository is MATLAB-focused. A separate portfolio draft mentions ROS/TurtleBot 3. Use the general MPC/robotics description unless a dedicated ROS repository or report is supplied.
- For the financial-markets project, use the repository-backed ARIMA description and approximately 82% accuracy. Do not use the separate 7% reinforcement-learning claim unless a supporting source is added.
- Do not fabricate client names, screenshots, testimonials, certifications, project metrics, or company-specific responsibilities.

6. Links and source provenance

LinkedIn: <https://www.linkedin.com/in/advait-jishnani>

GitHub profile: <https://github.com/advaitnaresh>

AI ATS repository: <https://github.com/advaitnaresh/shivarkATS>

MPC repository: <https://github.com/advaitnaresh/MPC>

Financial Markets repository: <https://github.com/advaitnaresh/Machine-Learning-In-Financial-Markets>

Restaurant Concierge repository: <https://github.com/advaitnaresh/Restaurant-Concierge-Bot>

Return on Creators: <https://returnoncreators.com/>

NYU Tandon Digital Learning: <https://engineering.nyu.edu/academics/programs/digital-learning>

Visa - About Visa: <https://corporate.visa.com/en/about-visa.html>

PERSONAL-SOURCE HIERARCHY

Primary personal source: the current master resume created from the agreed sb2nov / Minerva-style template. Secondary personal sources: prior resume variants, the current portfolio draft, and public repositories on Advait Jishnani's GitHub profile. Organization descriptions are concise paraphrases of official organization websites listed above.

DATA FRESHNESS

Prepared October 2026. Update this document whenever employment status, project links, project metrics, education details, or public contact information changes. The website should be regenerated from the newest version of this PDF, not from older resume variants.

END OF SOURCE-OF-TRUTH DOCUMENT